

# Cloud Computing for Science

## Part 3. Scaling Computation

Dennis Gannon

# The Cloud Data Center Evolution

- Early days: 2005
  - Very simple servers
  - Network outward facing poor interconnect
- 2008-2016
  - Software defined networks
  - Special InfiniBand sub networks
  - Many different server types
    - 2 cores to 32 cores to GPU accelerations
  - Efficiency experiments
    - Geothermal, wind, wave
    - Containerized server
- 2017
  - Azure FPGA accelerated mesh
  - Google Tensor Processing Unit
  - Facebook – Open Compute Project
  - ARM based servers

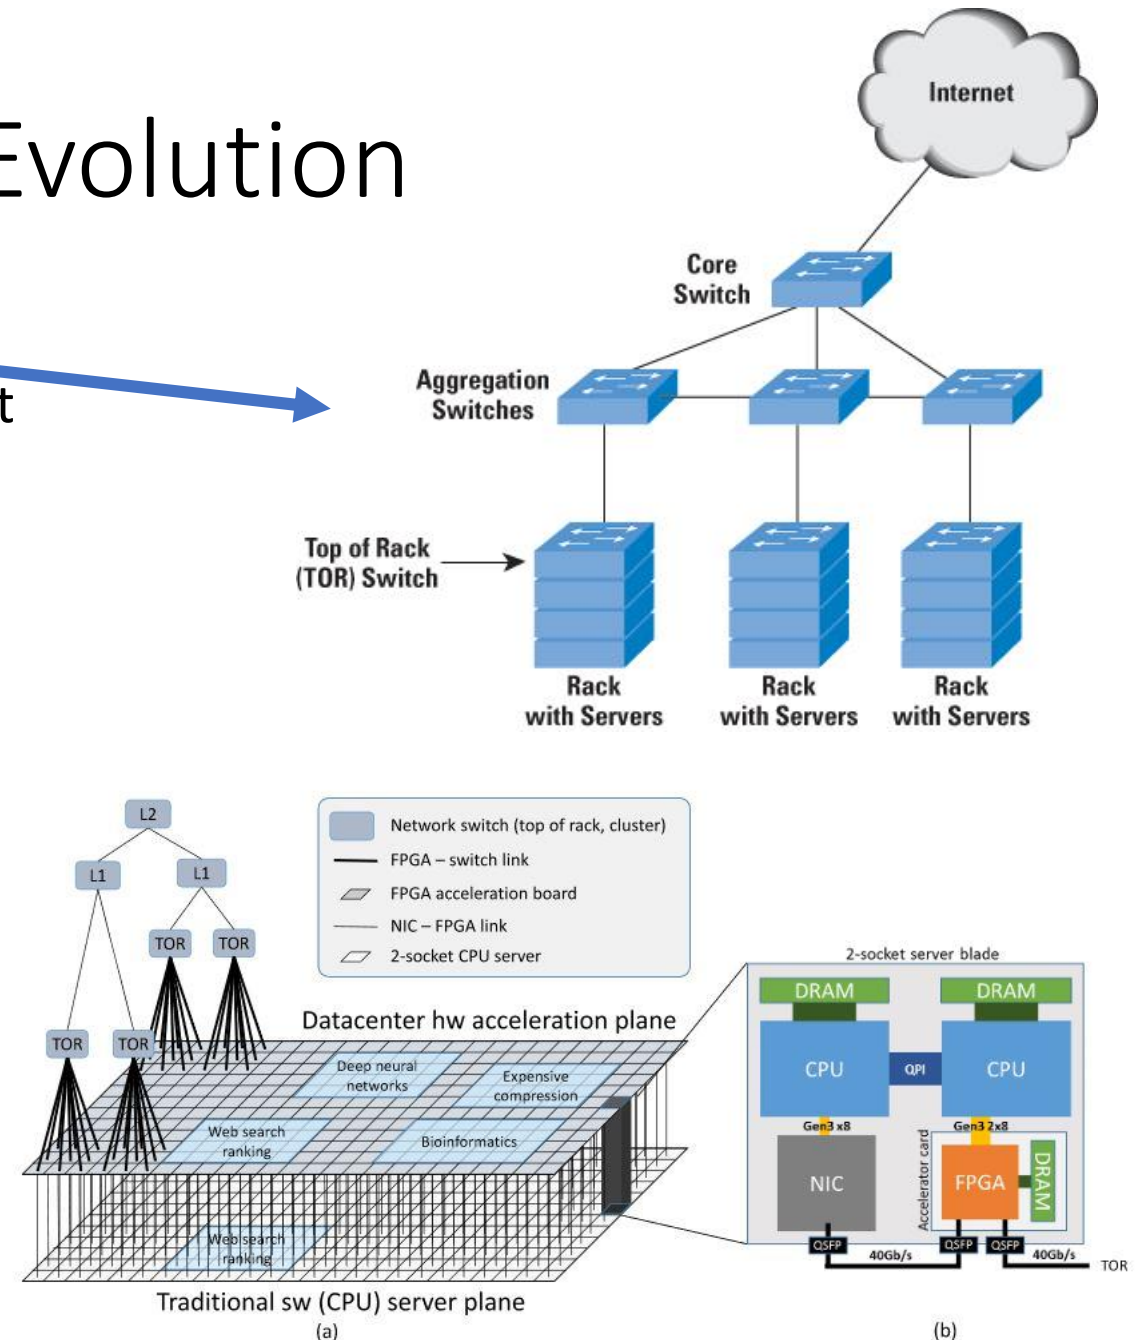


Fig. 1. (a) Decoupled Programmable Hardware Plane, (b) Server + FPGA schematic.

# Azure and AWS Global Data Center Network



# How to scale: Models of Parallelism

- Classic HPC
  - SPMD MPI programming
- Task Parallel
  - Also called “embarrassingly parallel”
- MapReduce
  - Hadoop style
- Graph Execution
  - Spark and streaming systems
- Microservices
  - Similar to actor model

# Classic HPC

- AWS CloudFormation Cluster
  - Fill out CfnCluster template
  - Use aws command line to submit
  - Log into head node
- Azure create a slurm cluster
  - See Azure slurm tutorial

## Deploy a slurm cluster

Deploy to Azure

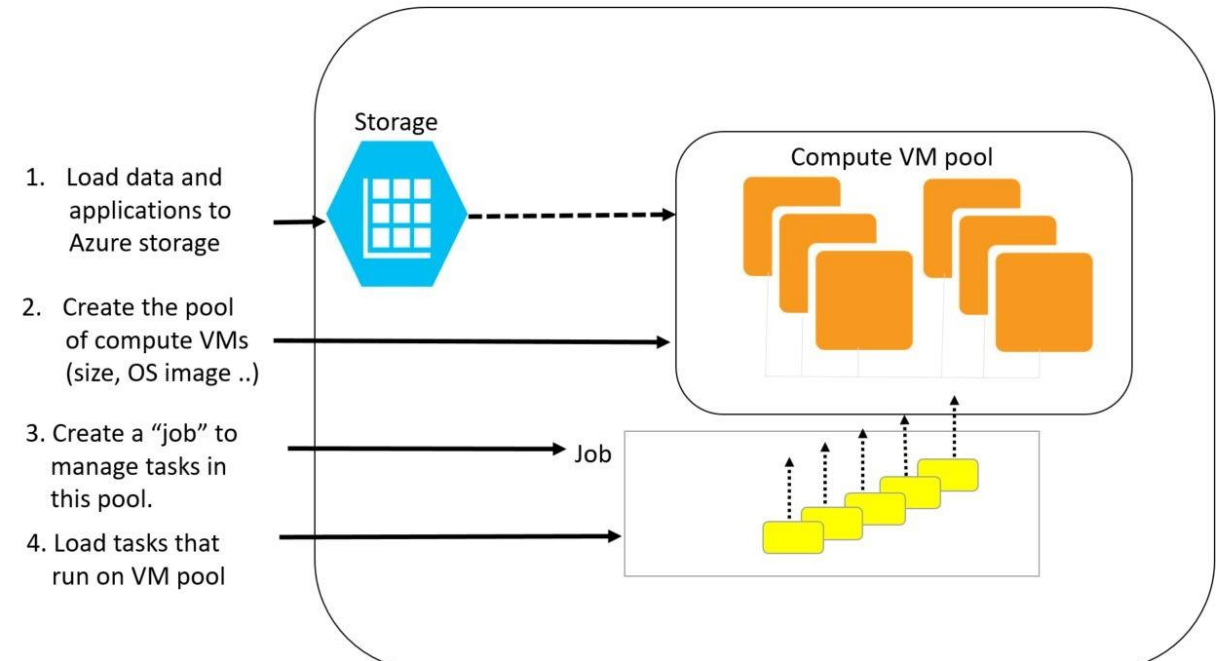
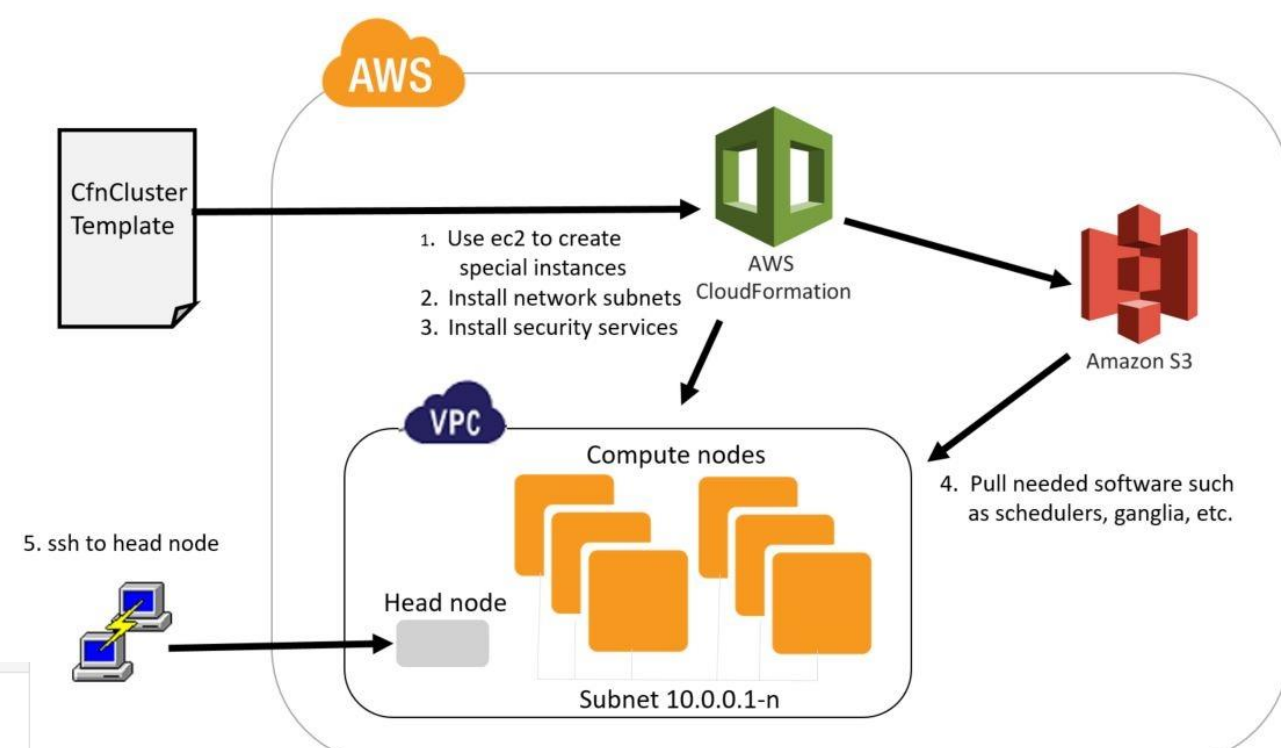
Visualize

1. Fill in the 3 mandatory parameters - public DNS name, a storage account to hold VM image, and admin user password.
2. Fill in other info and click "OK".

## Using the cluster

Simply SSH to the master node and do a `srun`! The DNS name is `dnsName.location.cloudapp.azure.com`, for example, `yidingslurm.westus.cloudapp.azure.com`.

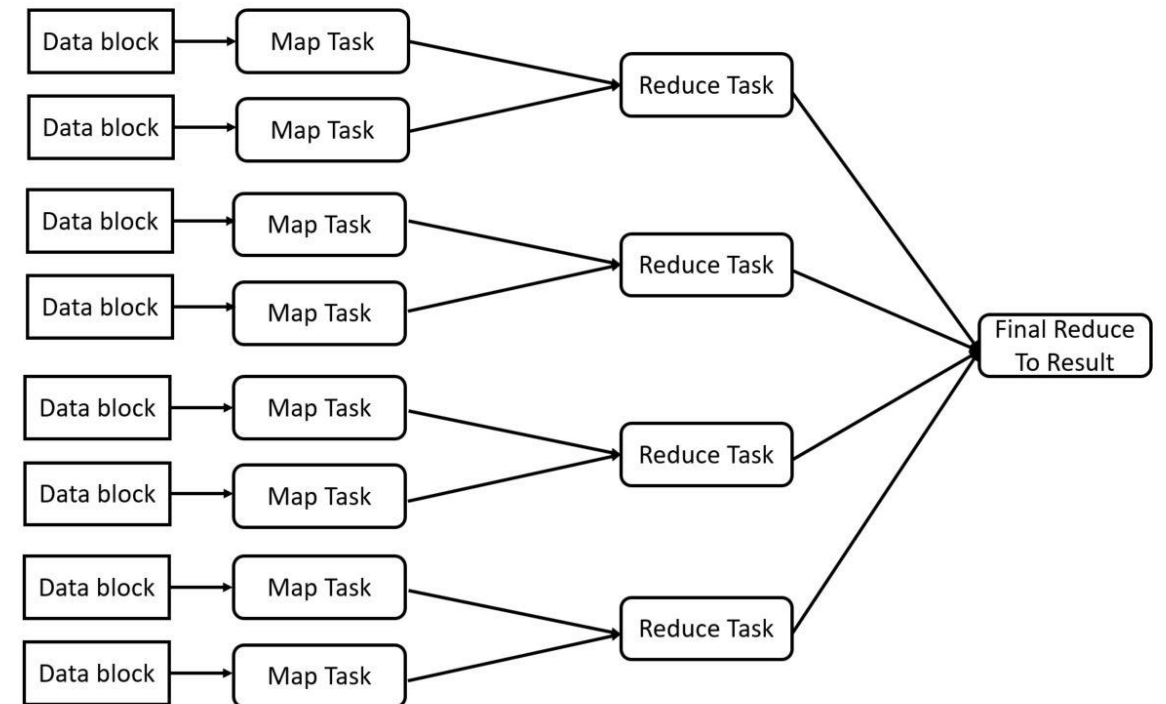
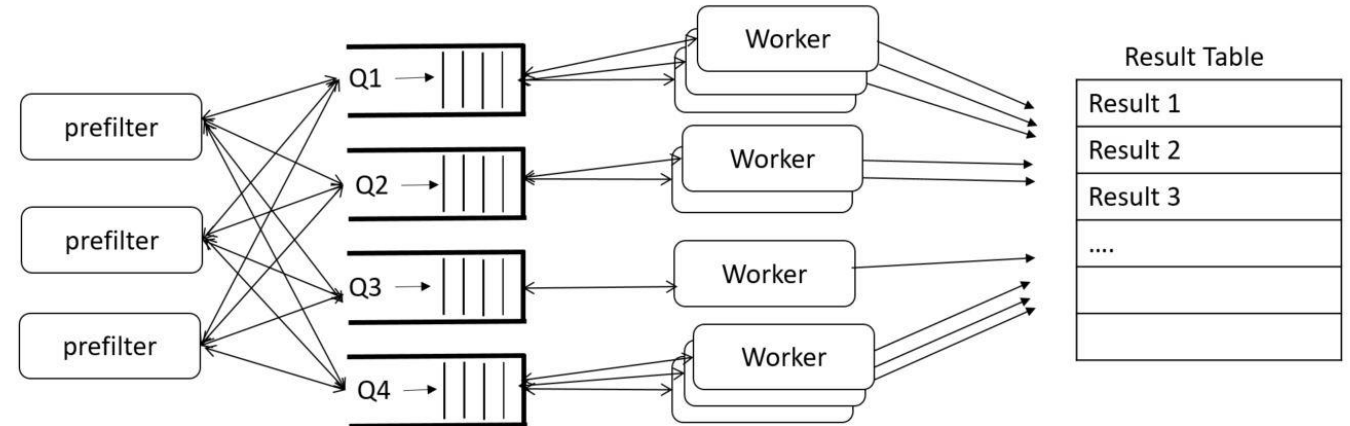
- Or use Azure Batch
  - Similar to AWS batch





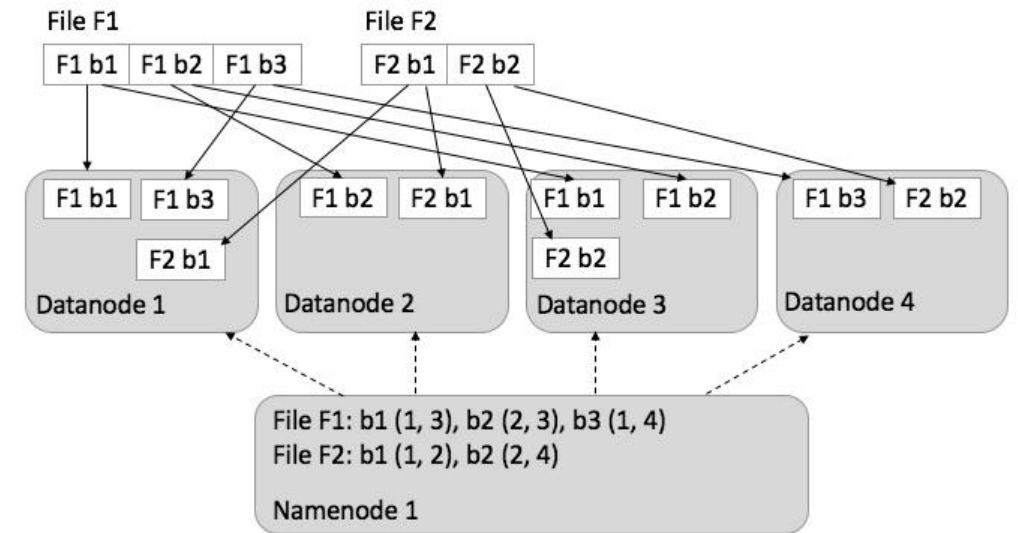
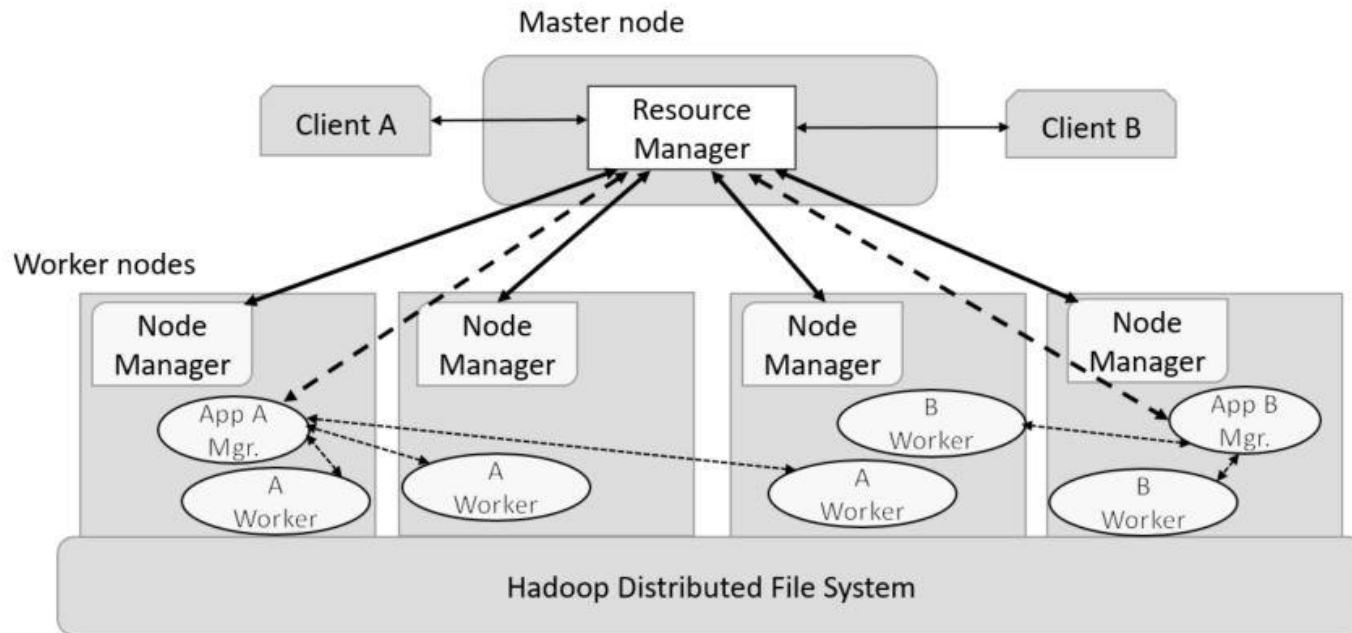
# Task Parallel and Map Reduce

- Task parallel model is great for solving problems that involve doing many independent computations.
- Map Reduce
  - Bulk Synchronous Parallel (BSP)
  - Map Task = an operation applied to blocks of data in parallel
  - Reduce Task- when maps are “done” reduce the results to a single result
- Examples of both later



# The Hadoop- Yarn ecosystem

- Yarn is the name of a project containing many elements
  - The runtime system is distributed
  - Hadoop, Spark run in distributed mode
  - Multiple clients can access the resource manager
  - Jupyter and Zeppelin are interactive clients
- HDFS is the Hadoop File system
  - Distributed over data node servers
  - Files are blocked, distributed and replicated
  - Files are write-once.



# Azure HDInsight is a Yarm Environment

Essentials ^

Resource group [\(change\)](#)  
[bookRG](#)

Status  
Running

Location  
East US

Subscription name [\(change\)](#)  
[azure4research](#)

Subscription ID  
f518fe6b-5262-4e5a-80cb-05b7a39f9298


Cluster type, HDInsight version  
Spark 2.0 on Linux (HDI 3.5.1000.0)


URL  
<https://bookcluster3.azurehdinsight.net>


Learn more  
[Documentation](#)  
[Getting Started](#)  
[Quickstart](#)

Head Nodes, Worker Nodes  
D3 (x2), D3 (x2)


Quick Links


 Cluster Dashboards

 Ambari Views

 Scale Cluster

Usage

  
HDInsight Cluster Dashboard

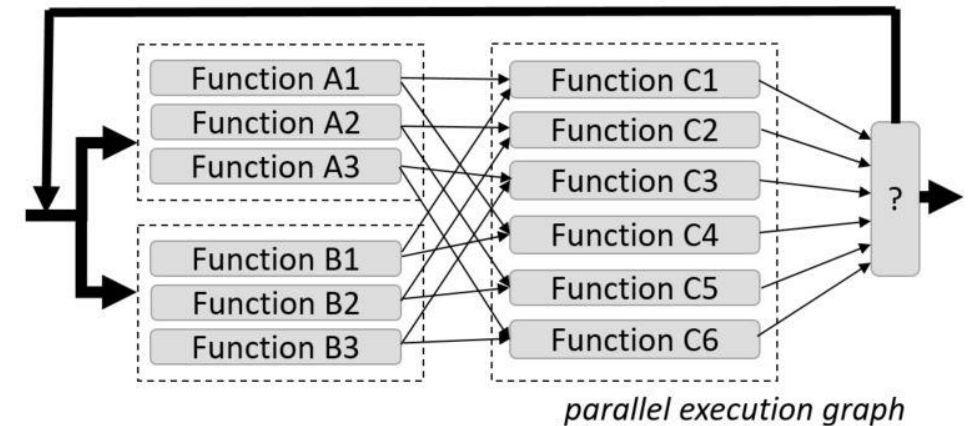
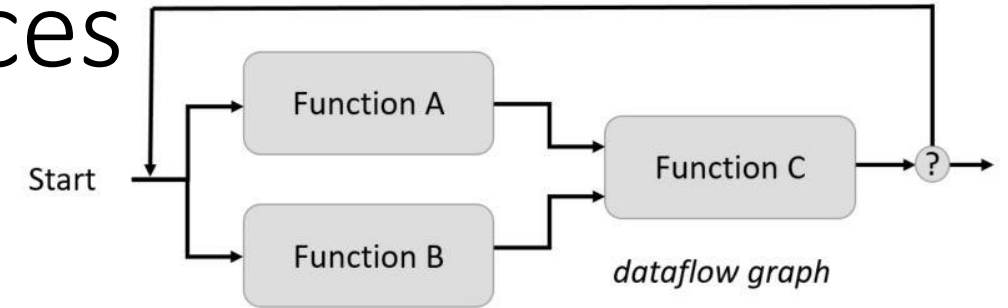
  
Jupyter Notebook



# Graph Parallel and Microservices

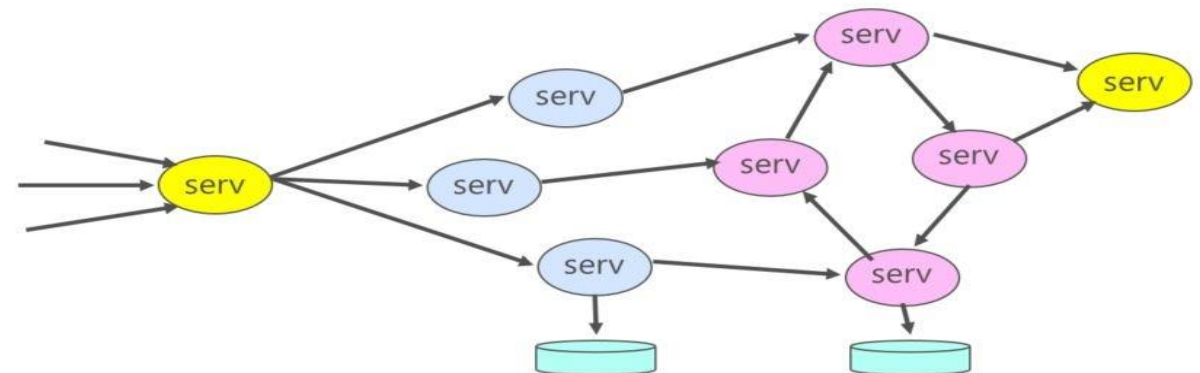
- Graph Parallel

- The data is in distributed arrays or streams.
- build a data flow graph of the algorithms functions.
- The graph is compiled into parallel operators that are applied to the distributed data structures.



- Microservices

- Divide a computation into small, mostly stateless components that can be
  - Easily replicated for scale
  - Communicate with simple protocols
- Computation is as a swarm of communicating workers.



# Graph computation example: Spark

- A simple map reduce: Compute
- For  $n = 10,000,000$
- In Spark on Python is:

$$\lim_{n \rightarrow \infty} \sum_{i=1}^n \frac{1}{i^2} = \frac{\pi^2}{6}$$

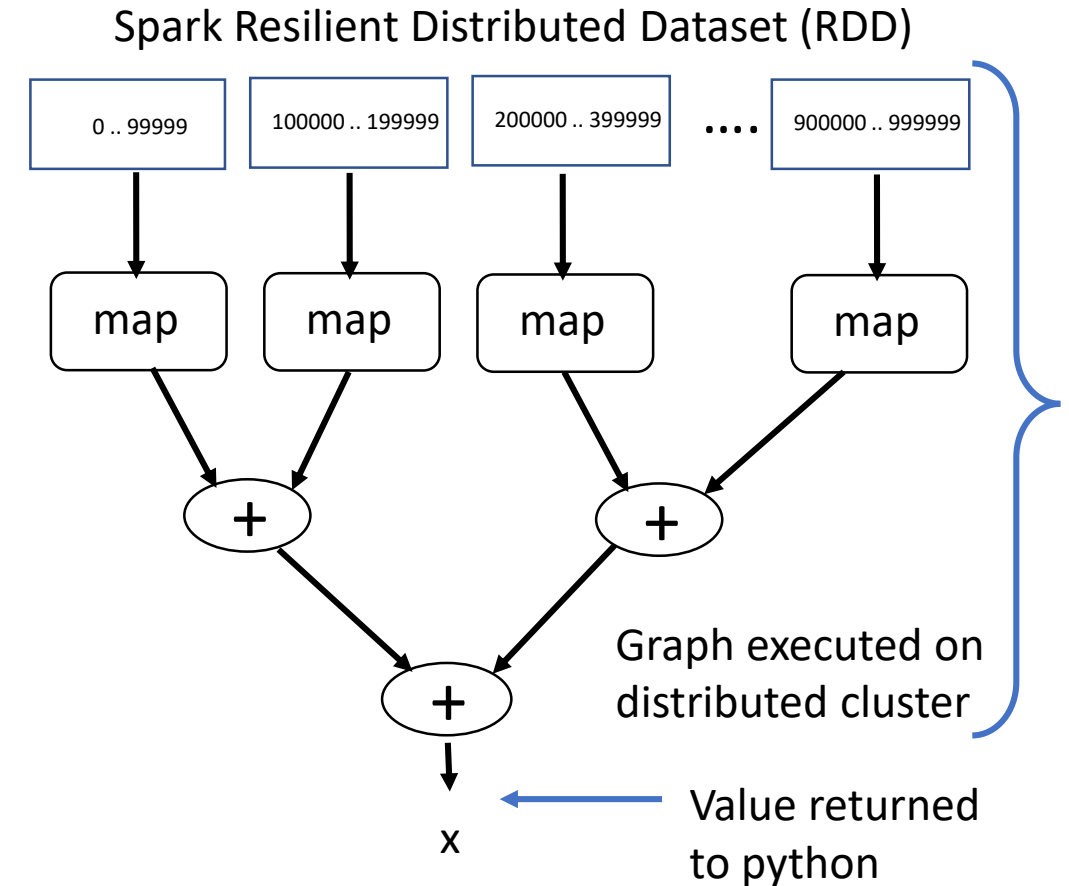
```
import numpy as np
ar = np.arange(n) # an array from 0 to 9999999
nump = 100

rdd = sc.parallelize(ar, nump)

x = rdd.map(lambda i: 1.0/(i+1)**2)
      .reduce(lambda a,b: a+b)

print("x=%f"%x)
print("pi**2/6=%f"%(np.pi**2/6))

1.644934
1.644934
```



# More interesting example: k-means clustering

- The algorithm basics
  - $n=1000000$
  - Start with a vector P of  $n$  2-d points and a vector kPoints of  $k$  random cluster centroids.
  - Iterate until kPoints don't move:
    - For each  $j$  in  $[0, k-1]$  pick  $q[j]$  from kPoints. Then find all the points  $p$  in P near  $q[j]$  and create the tuples  $(j, (p, 1))$  for  $p$  nearest to  $q[j]$
    - For each  $j$  compute the centroid of all points “near”  $q[j]$  in kPoints”  $(j, (\text{sum}(p)/\text{sum}(1)))$
    - Set  $q[j]$  to be the new centroid  $\text{sum}(p)/\text{sum}(1)$

```
tempDist = 1.0
while tempDist > convergeDist:
    newPoints = data \
        .map( lambda p: (closestPoint(p, kPoints), (p, 1))) \
        .reduceByKey(lambda x, y : (x[0] + y[0], x[1] + y[1])) \
        .map(lambda x : (x[0], x[1][0]/ x[1][1])) \
        .collect()

    tempDist = sum(np.sum((kPoints[i] - y) ** 2) \
                      for (i, y) in newPoints)
    for (i, y) in newPoints:
        kPoints[i] = y
```

# Section Summary

- The cloud data centers are designed to scale
  - Traditional HPC MPI programming is possible, but a Cray is better.
- The cloud is best at distributed scale, interactive computation
  - Spark in Yarn with Jupyter is a good example
- MapReduce and Graph models are well supported