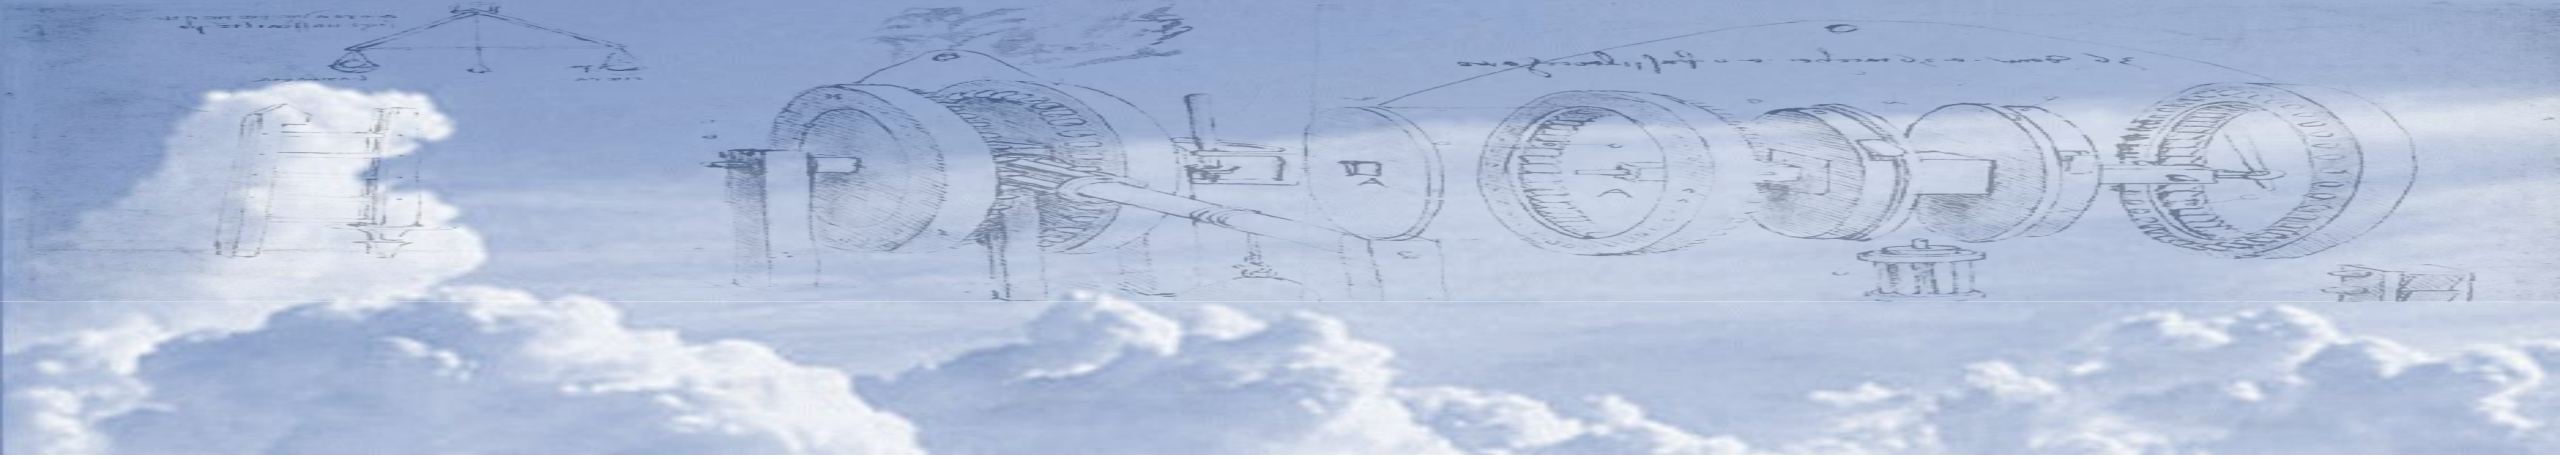

Cloud Computing for Science

Part 1. Managing Data in the Cloud

Dennis Gannon

Motivating Examples

- BIG objects
  - Climate scientists with big simulation output files in NetCDF.
    - Assume 10 TB in big objects
- Many small CVS files
  - Environmental Engineers with 1 Million records of observations each in CSV format
    - May be 100 TB total
- Streams
  - Scientists with a distributed collection of several thousand instruments.
    - Each generates a stream of records that must be collected and analyzed every few hours or continuously

Types of Cloud Data Storage Systems

- Basic Blob Object store
  - Buckets of immutable objects
  - Highly scalable & reliable
- Databases
  - SQL Style relational databases
  - NoSQL storage
  - Data warehouses
- Attached File stores
- Graph databases
- Streaming systems

Object stores

- Amazon AWS
  - S3- buckets of immutable objects
    - Organized in a 2-level folder system
    - Each object has associated metadata
- Microsoft Azure
  - Storage accounts contain blob storage along with tables, queues and file shares.
  - Blob containers are similar to S3
- Google Cloud Storage
  - Different models based on availably and cost
- OpenStack does not have a standard but semi-standards exist and various ones are used in various deployment

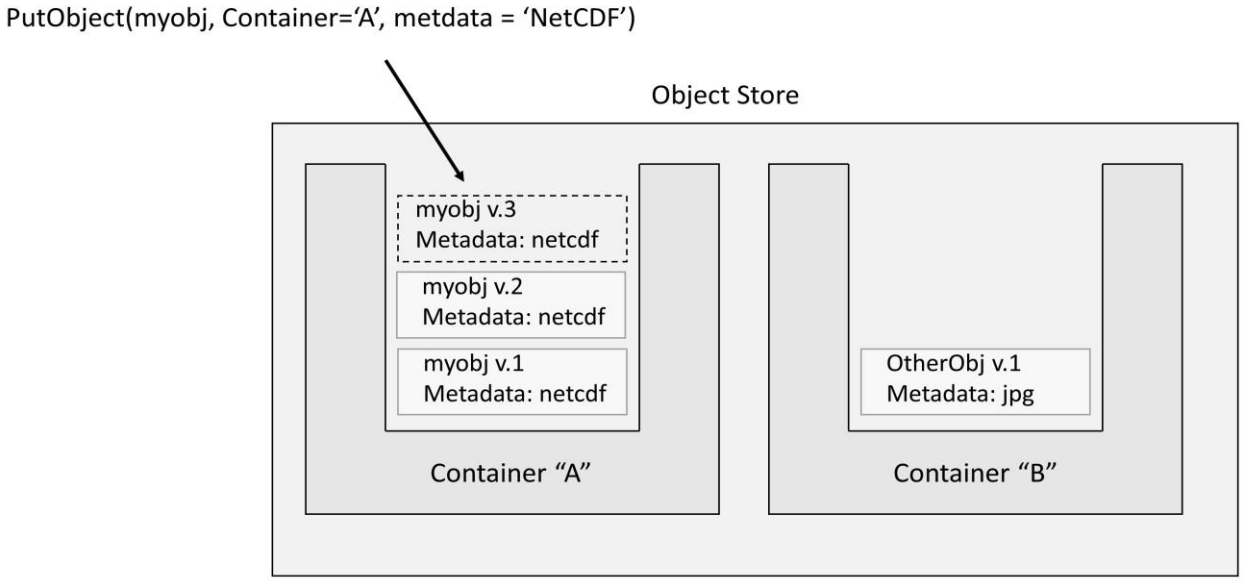

Azure storage account structure

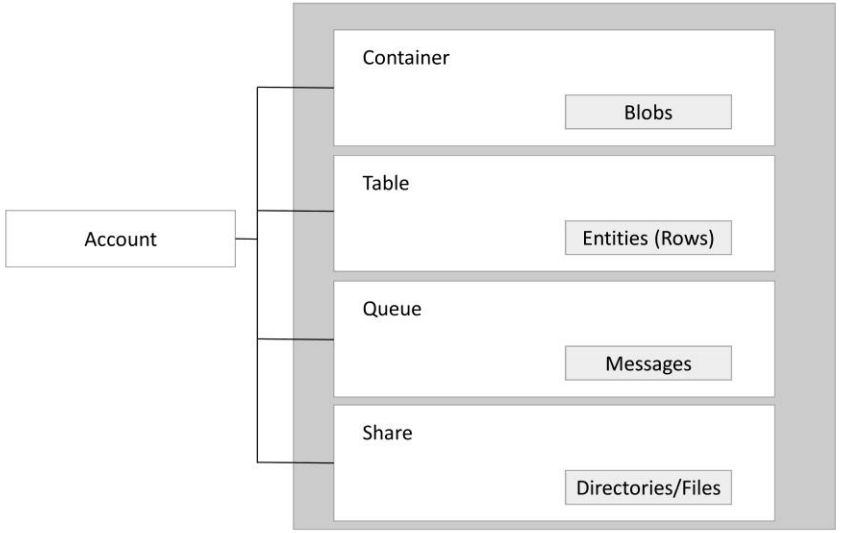

Relational Databases

- Amazon AWS
  - Relational Data Services (RDS)
  - Aurora
- Microsoft Azure
  - Azure SQL
    - 3 service tiers
    - Premium tier up to 1TB
    - Up to 30000 concurrent sessions
  - Azure Data Lake
- Google Cloud Storage
  - Cloud Spanner (beta)
    - Full relational
    - Strongly consistent
    - Scalable to thousands of servers

- Aurora is distributed
  - Scalable from 2 vCPUs to 32 vCPUs
  - Data up to 64TB
  - MySQL compatible
  - Fully geo-replicated

- Data Lake is a platform
  - Structured & unstructured data
  - Scalable to petabytes
  - Distributed and designed to support analytics

NoSQL Storage Systems

- Main Concepts
  - Designed for massive scale
    - Distributed over many storage nodes
  - Support some SQL operations (not joins)
  - May be only eventually consistent
  - Different types
    - Key-value
    - Column oriented
    - Document style
    - Graph

Relational table

Name	Job	Salary	age
Bill	Sales	$20000	18
Beth	Research	$1m	35
Carl	CEO	0	93
Jill	Prof	$100000	24

NoSQL Key-value system

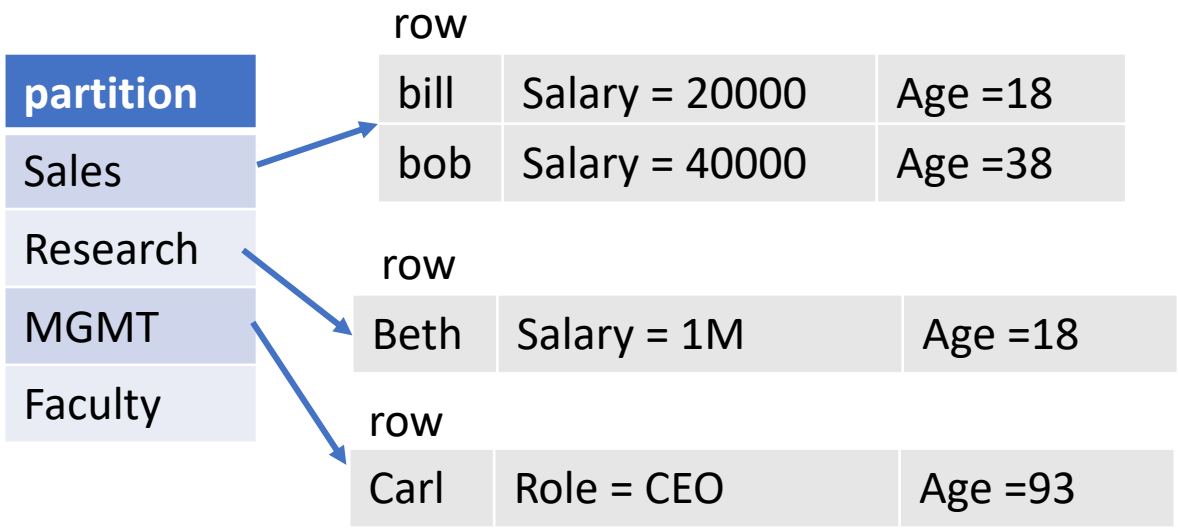

Key-value NoSQL examples *similar* to
Amazon AWS DynamoDB
Azure Tables
Google BigTable and Cloud DataStore

A simple example

- A table and data blobs
- Suppose you have a set of experiment data containing
  - An experiment number
  - A data item number
  - A date
  - Some comment data
  - A very large binary object
- Build a table of the experiments with a url link to the data in blob store.

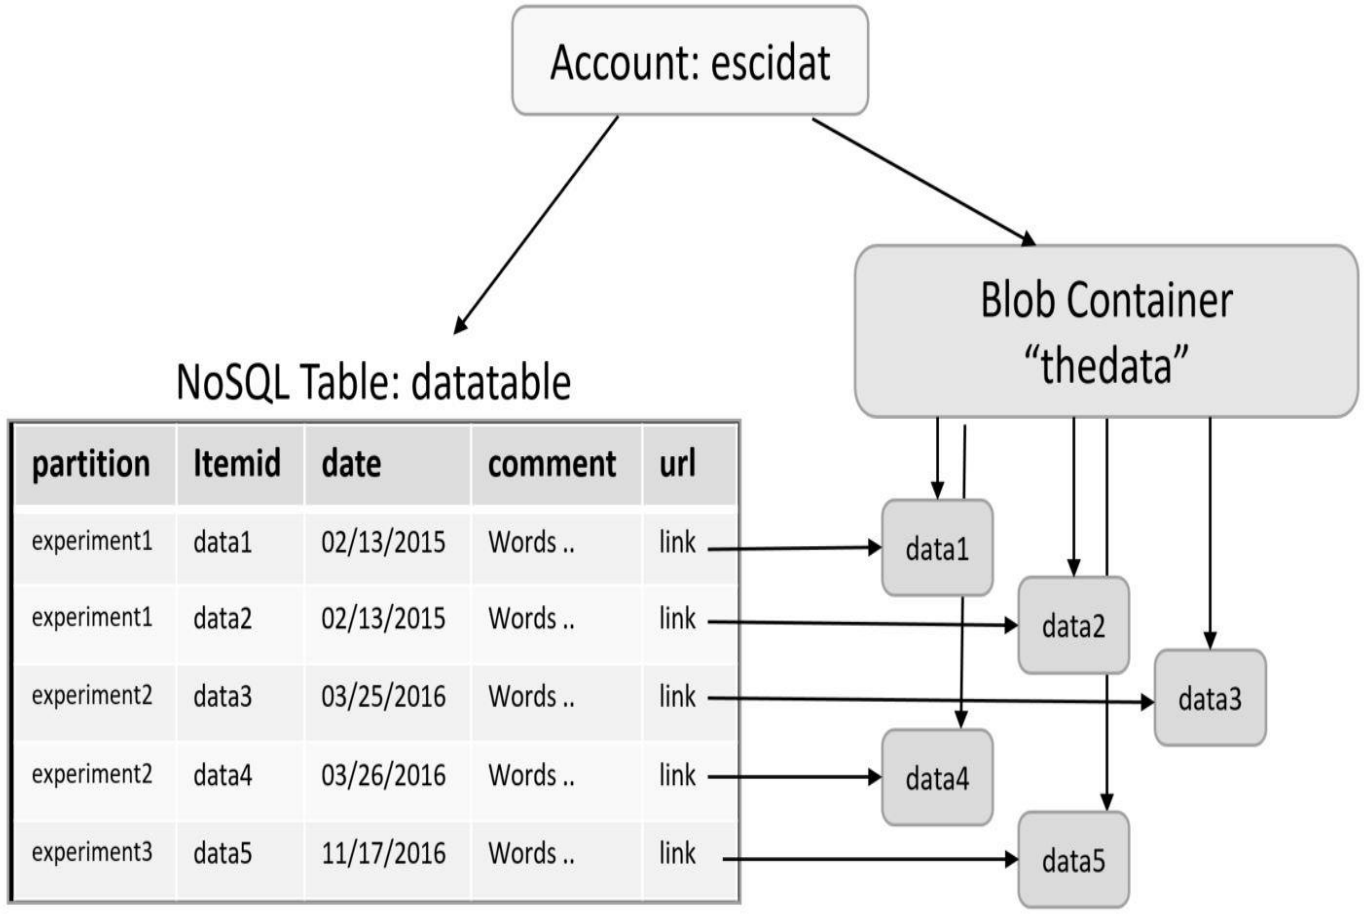

Azure Solution

Let’s create a new storage account “tutorial” to hold the blobs and table

To get an access key
Click here

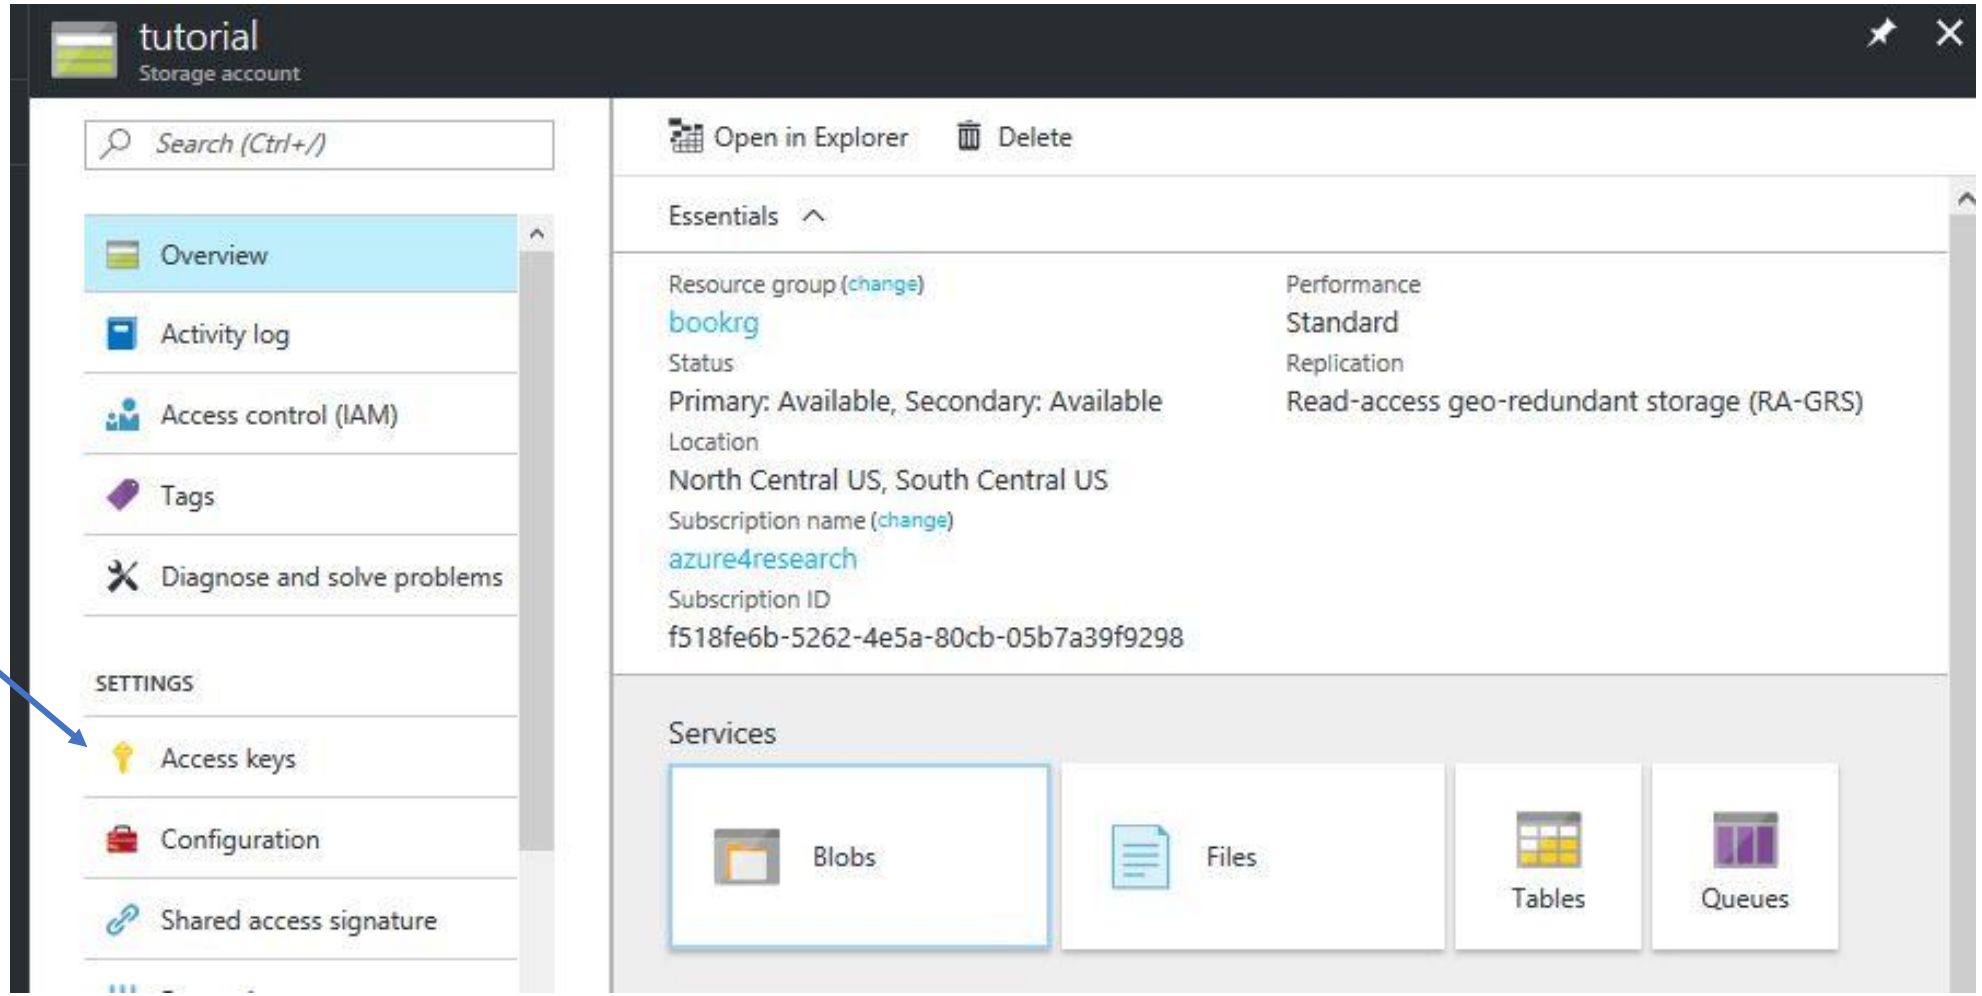

Get the container for the demo

- Make sure you have docker installed on your machine
- Download Docker for your pc or mac
  - https://docs.docker.com/engine/installation/
- Then do

  docker run -i -t -p 8888:8888 dbgannon/tutorial
- This will take a while
- When it is up go to https://localhost:8888
  - You will need to add security exceptions in the browser. It is safe.
  - Password is “tutorial”
  - Open azure.ipynb in Jupyter
- Or, if using a different jupyter, download https://SciengCloud.github.io/azure.ipynb
- Using the azure portal create a storage account and have the key ready

Azure Solution – now create table and blob container

```
import csv
import sys
import azure.storage
from azure.storage.table import TableService, Entity
from azure.storage.blob import BlockBlobService
from azure.storage.blob import PublicAccess

block_blob_service = BlockBlobService(account_name='tutorial',
      account_key='biglongaccesskey')

block_blob_service.create_container('datacont',
            public_access=PublicAccess.Container)

table_service = TableService(account_name='tutorial',
      account_key='samebiglongkey')


if table_service.create_table('DataTable'):
    print "table created"
else:
    print "table already there"
```

Azure Solution

- Assume data objects are stored in files in /home/me/ and there is a CSV file “thedata” with rows
  - Experiment name, item id, date, filename, comment string

```
with open('/datadir/experiments.csv', 'rb') as csvfile:
    csvf = csv.reader(csvfile, delimiter=',', quotechar='|')
    for item in csvf:
            print item
            block_blob_service.create_blob_from_path(
                    'datacont', item[3],
                   "/datadir/"+item[3]
             )
        url = "https://"+account+".blob.core.windows.net/datacont/"+item[3]
        metadata_item = {'PartitionKey': item[0], 'RowKey': item[1],
                         'description' : item[4], 'date' : item[2], 'url':url}
        table_service.insert_entity('DataTable', metadata_item)
```

Use Azure Storage Explorer to inspect the table

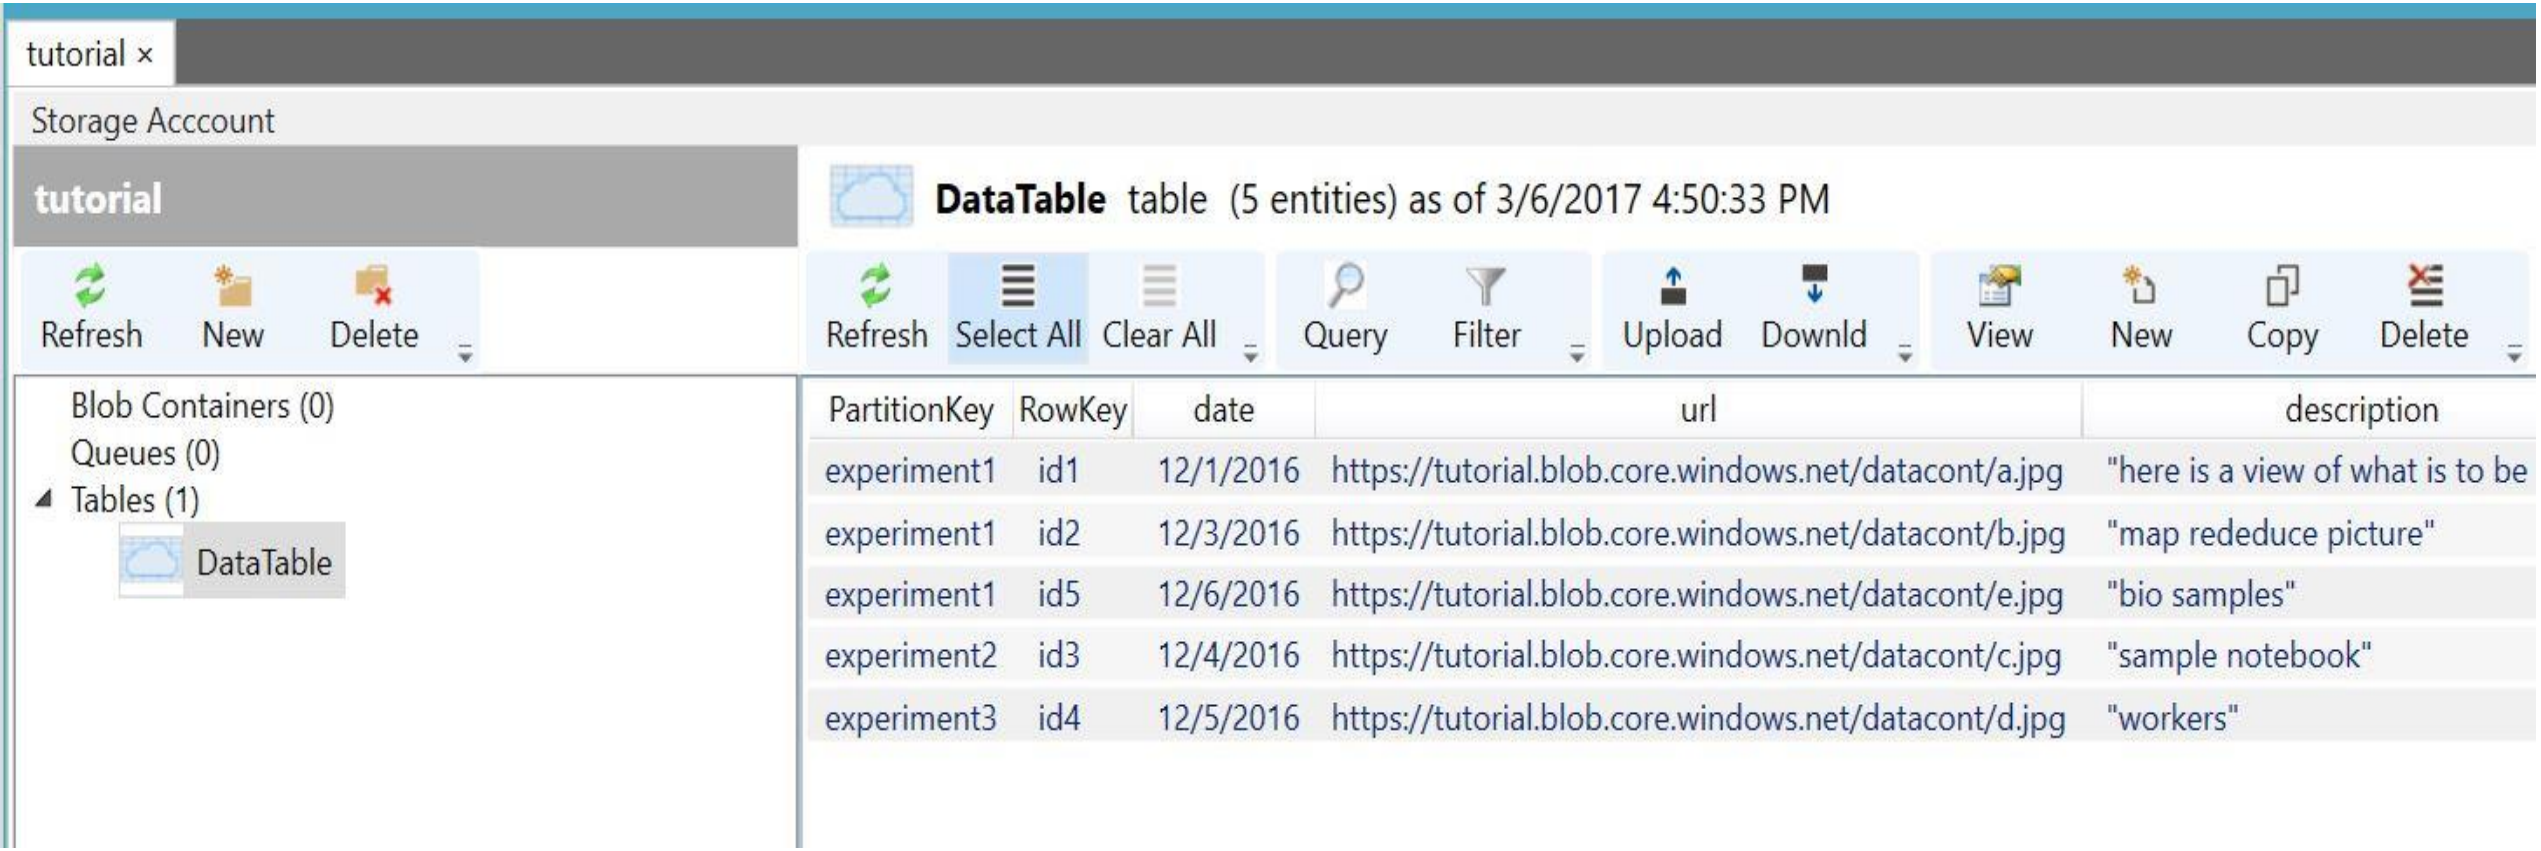

Next VMs and Containers